

# CS140E: embedded OS

Stanford win 2025

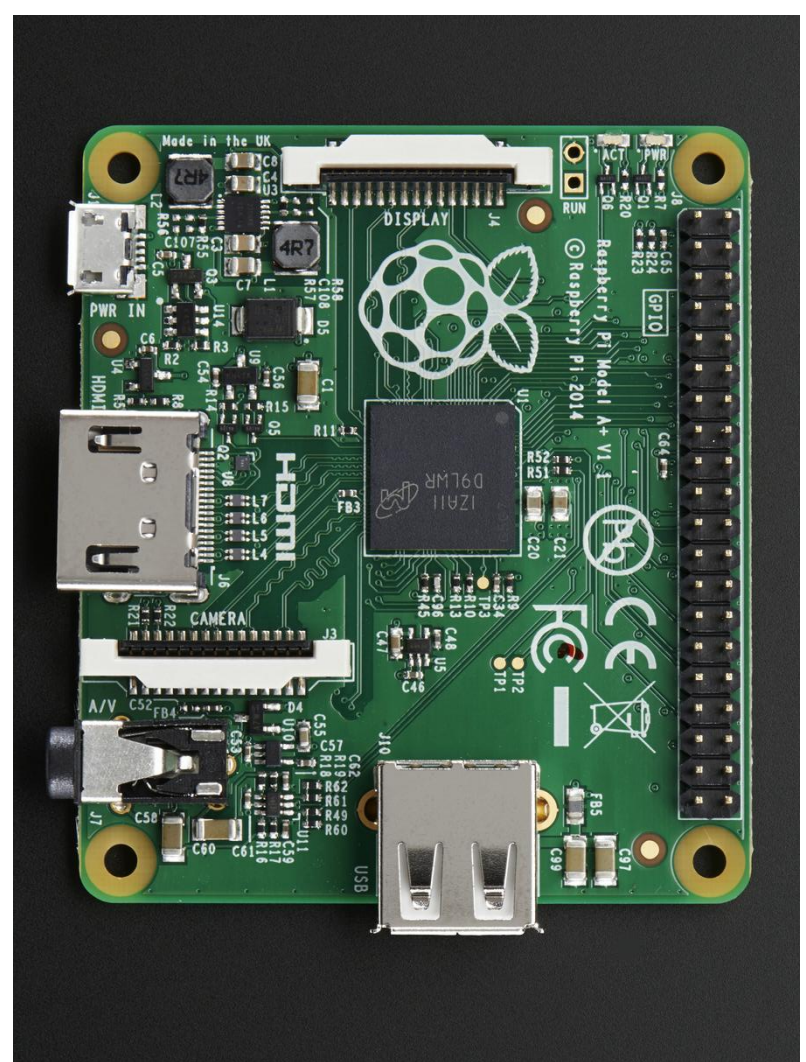
# Outline

- Staff:
  - Dawson Engler (“Dawson” is fine: if stuff is broken it’s my fault)
  - CAs: Joe Tan, Matthew Sotoudeh (paid for outside class stuff)
  - SLs: Joseph Shetaye, Arjun Vikram (not paid for outside class stuff)
  - Volunteers: Ammar Ali Ratnani, Aaryan Singhal (not paid at all 🤖)
- What
- Why

# What

Write small, clean OS on an r/pi A+:

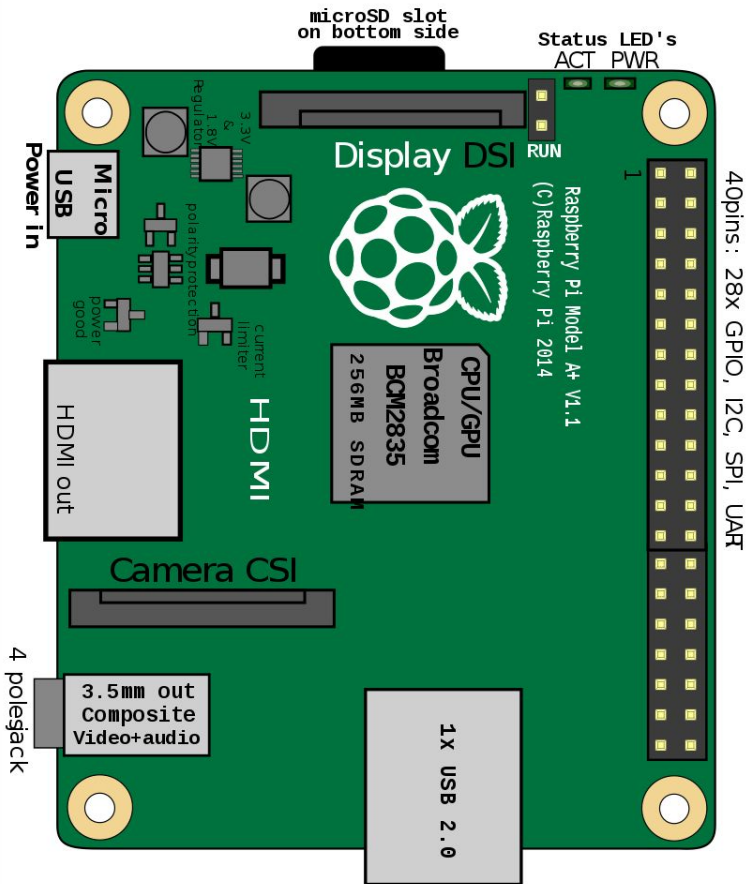
[https://en.wikipedia.org/wiki/Raspberry\\_Pi](https://en.wikipedia.org/wiki/Raspberry_Pi)



# What

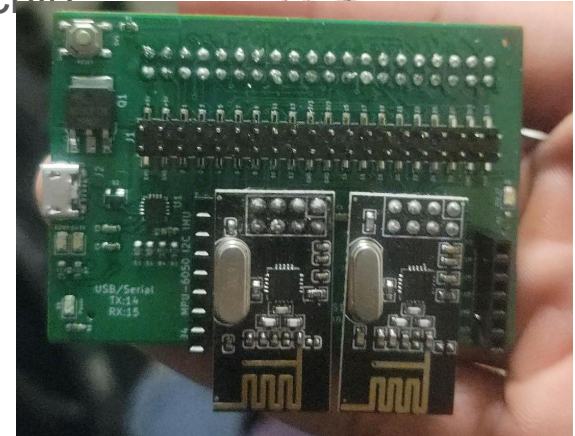
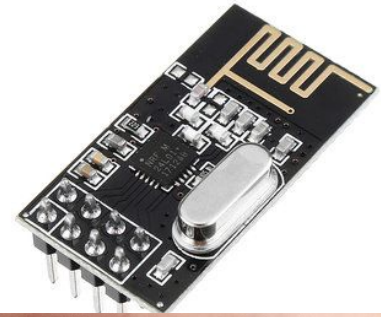
Write small, clean OS on an r/pi A+:

[https://en.wikipedia.org/wiki/Raspberry\\_Pi](https://en.wikipedia.org/wiki/Raspberry_Pi)



# What you will build (different than last year)

- Bootloader (ship code from your laptop to the pi)
- GPIO, UART, SPI Device drivers
- Threads + interrupts
- Virtual Memory
- Simple fat32 file system on SD card.
- Networking code that uses the cheap nrf24L01 chip
  - Will use a pcb that Parthiv designed as 340lx project
- End result: simple, small unix OS
  - Mostly your code



# Why OS?

If you can write a real OS, you can write almost anything (non-math-y).

Once you get this, easy to delta to something else.

Classes are fake: real world is not a clean, textbook of systematized knowledge

Difficult to understand documents

Wrong

Incomplete

Not written to be used

You will learn to operate in such a world without a lot of panic/drama.

# Your OS

You will write (almost) all code.

Small + lightweight = you can do things impossible on modern OSes.

Nanosecond latencies for messages

Real time guarantees faster than expensive digital tools

Exception tricks that let you build valgrind in < 1KLOC.

....

# Why R/PI A+?

Most OSes write code on a fake simulator

Alot of work, not that cool at the end.

R/pi = real computer for about \$20 and an ounce of weight.

Many examples / blog posts of how to do various things.

Unlike most machines, makes interacting with the real world easy.

Can build many interesting systems b/c can use weird hardware easily  
(motion sensors, IR sensors, accelerometer, gyroscope, light sensor...)



# Class philosophy

Write a complete, narrow OS all the way down to the bare metal.

We cover less material than most OS classes (multi-level schedulers, multi-level page tables)

However, you will understand much more thoroughly

Hope: easy to do delta off of your knowledge to more fancy things.

Labs vs lectures:

Always try to have you be writing code. You will actually understand what is going on.

Common tragedy of OS: missing a key sentence, mistake in key document. We will use lab to fill this in, saving you many hours/days.

Goal: you do pre-work to pre for lab, walk in, by the end of the lab, you have a complete working simple version of a key trick.

# Class philosophy, subtext

- This class is for people that are super interested in low level
  - If you're on the spectrum (hi): this class is designed for you.
  - Our goal: move fast with a bunch of people that are very interested.
  - Only take the class if (1) stanford is easy or (2) very interested and have an open quarter.
  - Our hope: we don't waste your time, and you cover adult stuff in a way and level haven't seen.
  - One result: you'll find the hardcore people at stanford you want to work with in the future.
- Warning:
  - Don't take this class thinking it is an "easier version of 112"
  - It is not.
  - Making harder: we don't have enough staff!

# Goal: you will develop two super-powers

Power 1: Differential debugging.

Efficiently answering “why doesn’t work” for complex things.

Swap working pieces + Binary search

Power 2: Epsilon development.

Epsilon sprint paradox: When building systems, the smaller the step you take, the faster you can run.

# Differential debugging

You write code, it doesn't work, the error could be:

- The code you wrote

- Hardware fault (bad manufacturing, smoked something)

- Wiring mistake

- Subtle cache issue

- Compiler problem (more on this)

- ...

You will get good at breaking down problems by swapping pieces between a working system (yesterday's code, your partners lab) and a non-working system (today's code, your lab)

# Example from next lab.

You get the following set of stuff:

To run you:

Copy blink.bin to sd

Wire up led

Wire up serial device

Plug into your laptop

It doesn't work.

But your partner's does



# Example from next lab.

You get the following set of stuff:

To run you:

Copy blink.bin to sd

Wire up led

Wire up serial device

Plug into your laptop

It doesn't work.

But your partner's does



# Example from next lab.

You get the following set of stuff:

To run you:

Copy blink.bin to sd

Wire up led

Wire up serial device

Plug into your laptop

It doesn't work.

But your partner's does



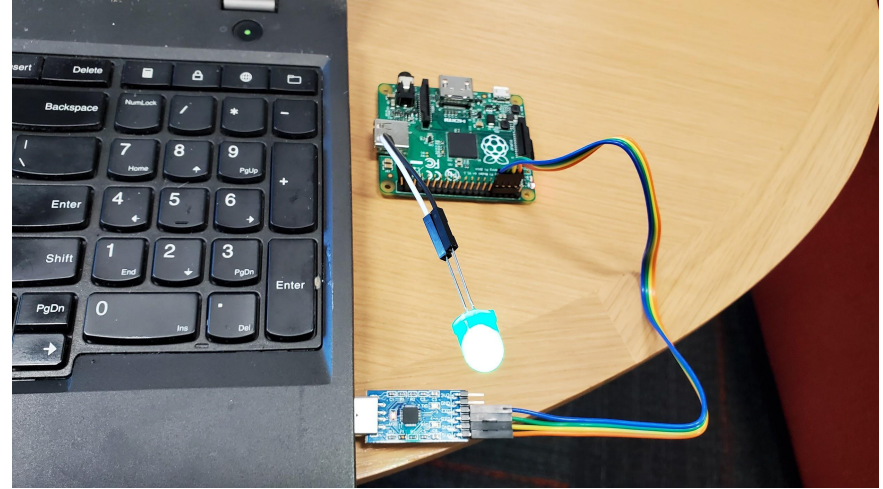


# What is messed up?

Yours: Not working



Partner's: Working




What to do first?




# What is messed up?

Yours: Not working



Partner's: Working




What does swapping tell you if doesn't work?


What does swapping tell you if works?

# What is messed up?

Yours: Not working



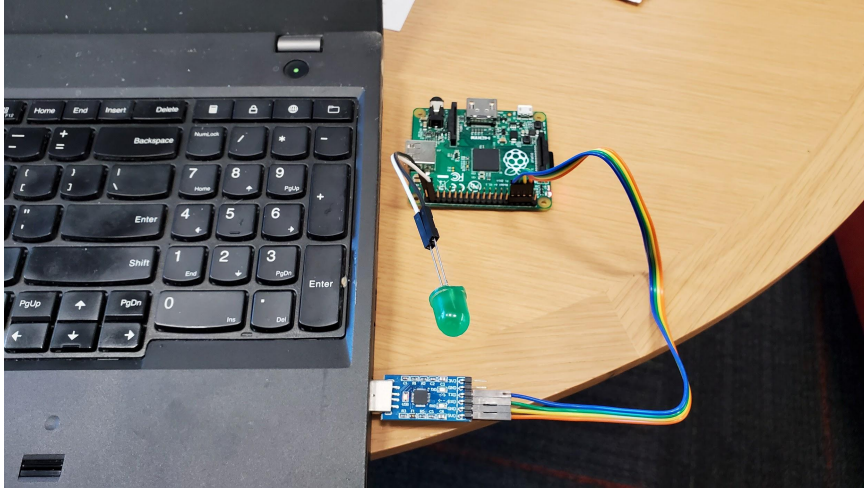
Partner's: Working



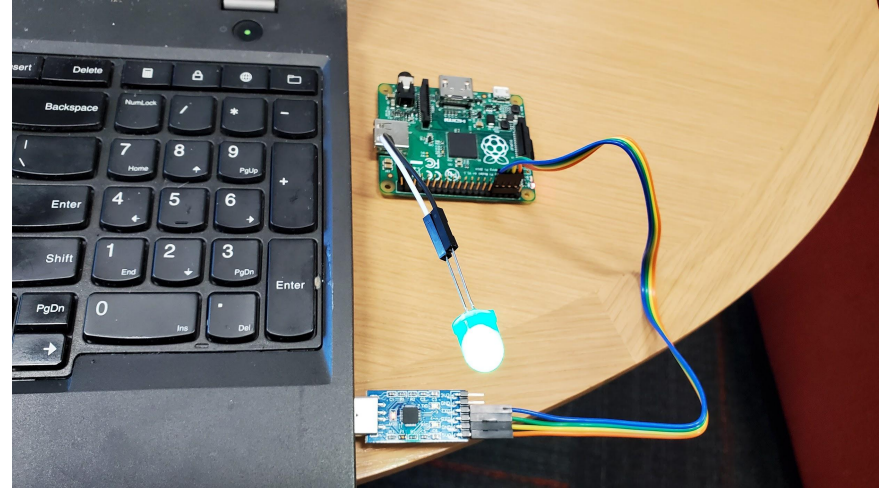
Swapping works: how to narrow down with least work?

# What is messed up?

Yours: Not working



Partner's: Working



Entire class: whenever we control device, has some software component  $S$  (can be wrong) and some hardware component  $H$  (can be broken).

Doesn't work = linear equation solving with two variables. How to isolate?

# Epsilon sprinting: Slow is fast.

What is wrong?

If I did X, it's X.

If I did  $X_1 + X_2 + \dots + X_n$  it could be any, or some combination.

Inverting crash / bug to root cause is much harder in the latter case.

My epsilon-sprint theorem:

Given a working system  $W$  and a change  $C$ , then as  $|C| \Rightarrow \epsilon$ , the time + computation (IQ) it takes to figure out why  $\{ W + C \}$  doesn't work goes to 0.

Related claim: the time it takes to debug why a change broke the system increases non-linearly with the size of the change.

# Administrivia

- Rules:
  - Feel free to leave early, but don't be late (makes us run  $O(n)$ )
  - Don't use co-pilot or other people's code. You have to write.
  - Don't miss more than 2 labs (even that is very tough).
- Grade:
  - "A" requires 3+ hard extensions
  - Absolutely must turn in lab w/in 7 days. No exceptions!
  - Usually prelab due before lab. Won't accept after.
  - Will have a final project (roughly 3+ labs worth of work)
  - Will have 2 homeworks.
- Plus side:
  - We pay for food.
  - We pay for equipment
  - We stay til 11:30pm or so.

# Administrivia

Two labs each week.

Each lab will have pre-lab work you should turn in before lab.

Ideally finish during the lab period (I will stay til everyone is done).

Must finish within a week of the lab, or start losing a letter grade each day.

Must pre-arrange missed labs. It's a problem to miss more than a couple.

There (tentatively) will be three “capstone” homework assignments that consolidate a chunk of labs together.

If you've done the lab, this shouldn't be a big deal.

# Administrivia

You can work with other people!

However, you *\*must\** type and turn in everything yourself.

Please post to the newsgroup.

If the rules on in-person meetings lift, it's great if you can do the labs in groups.

(We will pay for food you order during lab if this happens.)



# What to do now

- clone the class git repository:
  - <https://github.com/dddrreee/cs140e-24win>
  - `git clone git@github.com:dddrreee/cs140e-24win.git`
- For lab next tues, make sure you:
  - have a way to write either a micro-SD or SD card.
  - 
  - 
  - Have a way to plug in a standard USB device
  - 
  - Do PRELAB for lab 2-trust

